

```
%matplotlib inline
```

Double-click (or enter) to edit

```
from google.colab import drive
drive.mount('/content/drive')
```

☞ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount(

```
import os
os.chdir('/content/drive/My Drive/cs505/char_rnn_tutorial') #change dir
!pwd
```

☞ /content/drive/My Drive/cs505/char_rnn_tutorial

Classifying Names with a Character-Level RNN

Author: Sean Robertson <<https://github.com/spro/practical-pytorch>>_

We will be building and training a basic character-level RNN to classify words. A character-level RNN is outputting a prediction and "hidden state" at each step, feeding its previous hidden state into each new step. The output, i.e. which class the word belongs to.

Specifically, we'll train on a few thousand surnames from 18 languages of origin, and predict which language the word belongs to.

::

```
$ python predict.py Hinton
(-0.47) Scottish
(-1.52) English
(-3.57) Irish
```

```
$ python predict.py Schmidhuber
(-0.19) German
(-2.48) Czech
(-2.68) Dutch
```

Recommended Reading:

I assume you have at least installed PyTorch, know Python, and understand Tensors:

- <http://pytorch.org/> For installation instructions
- :doc: /beginner/deep_learning_60min_blitz to get started with PyTorch in general
- :doc: /beginner/pytorch_with_examples for a wide and deep overview

- :doc: /beginner/former_torchies_tutorial if you are former Lua Torch user

It would also be useful to know about RNNs and how they work:

- The Unreasonable Effectiveness of Recurrent Neural Networks <<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>> __ shows a bunch of real life examples
- Understanding LSTM Networks <<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>> but also informative about RNNs in general

▼ Preparing the Data

.. Note:: Download the data from here <<https://download.pytorch.org/tutorial/data.zip>> _{
Included in the `data/names` directory are 18 text files named as "[Language].txt". Each file contains a mostly romanized (but we still need to convert from Unicode to ASCII).

We'll end up with a dictionary of lists of names per language, `{language: [names ...]}`. The general language and name in our case) are used for later extensibility.

```
from __future__ import unicode_literals, print_function, division
from io import open
import glob
import os

def findFiles(path): return glob.glob(path)

print(findFiles('data/cities_train/*.txt'))

import unicodedata
import string

all_letters = string.ascii_letters + " .,:;"
n_letters = len(all_letters)

# Turn a Unicode string to plain ASCII, thanks to http://stackoverflow.com/a/518232/2
def unicodeToAscii(s):
    return ''.join(
        c for c in unicodedata.normalize('NFD', s)
        if unicodedata.category(c) != 'Mn'
        and c in all_letters
    )

print(unicodeToAscii('Ślusàrski'))

# Build the category_lines dictionary, a list of names per language
category_lines = {}
val_category_lines = {}
all_categories = []
val_categories = []
```



```

    return all_letters.find(letter)

# Just for demonstration, turn a letter into a <1 x n_letters> Tensor
def letterToTensor(letter):
    tensor = torch.zeros(1, n_letters)
    tensor[0][letterToIndex(letter)] = 1
    return tensor

# Turn a line into a <line_length x 1 x n_letters>,
# or an array of one-hot letter vectors
def lineToTensor(line):
    tensor = torch.zeros(len(line), 1, n_letters)
    for li, letter in enumerate(line):
        tensor[li][0][letterToIndex(letter)] = 1
    return tensor

print(letterToTensor('J'))

print(lineToTensor('Jones').size())

[> tensor([[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
           0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
           0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
           0., 0., 0.]])
      torch.Size([5, 1, 57])

```

▼ Creating the Network

Before autograd, creating a recurrent neural network in Torch involved cloning the parameters of a lay hidden state and gradients which are now entirely handled by the graph itself. This means you can im regular feed-forward layers.

This RNN module (mostly copied from the PyTorch for Torch users tutorial <http://pytorch.org/tutorials/beginner/former_torchies/nn_tutorial.html#example-2> which operate on an input and hidden state, with a LogSoftmax layer after the output.

.. figure:: <https://i.imgur.com/Z2xbySQ.png> :alt:

```

import torch.nn as nn

class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(RNN, self).__init__()

        self.hidden_size = hidden_size

        self.i2h = nn.Linear(input_size + hidden_size, hidden_size)
        self.i2o = nn.Linear(input_size + hidden_size, output_size)
        self.softmax = nn.LogSoftmax(dim=1)

```

```

def forward(self, input, hidden):
    combined = torch.cat((input, hidden), 1)
    hidden = self.i2h(combined)
    output = self.i2o(combined)
    output = self.softmax(output)
    return output, hidden

def initHidden(self):
    return torch.zeros(1, self.hidden_size)

```

```

n_hidden = 128
rnn = RNN(n_letters, n_hidden, n_categories)

```

To run a step of this network we need to pass an input (in our case, the Tensor for the current letter) and initialize as zeros at first). We'll get back the output (probability of each language) and a next hidden state.

```

input = letterToTensor('A')
hidden = torch.zeros(1, n_hidden)

output, next_hidden = rnn(input, hidden)
print(output)

```

```

↳ tensor([[ -2.1942, -2.1523, -2.1505, -2.2985, -2.1375, -2.2217, -2.2496, -2.2357,
           -2.1476]], grad_fn=<LogSoftmaxBackward>)

```

For the sake of efficiency we don't want to be creating a new Tensor for every step, so we will use `lineToTensor` and use slices. This could be further optimized by pre-computing batches of Tensors.

```

input = lineToTensor('Albert')
hidden = torch.zeros(1, n_hidden)

output, next_hidden = rnn(input[0], hidden)
print(output)

```

```

↳ tensor([[ -2.1942, -2.1523, -2.1505, -2.2985, -2.1375, -2.2217, -2.2496, -2.2357,
           -2.1476]], grad_fn=<LogSoftmaxBackward>)

```

As you can see the output is a `<1 x n_categories>` Tensor, where every item is the likelihood of the

▼ Training

Preparing for Training

Before going into training we should make a few helper functions. The first is to interpret the output of the network as a likelihood of each category. We can use `Tensor.topk` to get the index of the greatest value:

```
def categoryFromOutput(output):
    top_n, top_i = output.topk(1)
    category_i = top_i[0].item()
    return all_categories[category_i], category_i
```

```
print(categoryFromOutput(output))
```

```
↳ ('ir', 4)
```

We will also want a quick way to get a training example (a name and its language):

```
import random
```

```
def randomChoice(l):
    return l[random.randint(0, len(l) - 1)]
```

```
def randomTrainingExample():
    category = randomChoice(all_categories)
    line = randomChoice(category_lines[category])
    category_tensor = torch.tensor([all_categories.index(category)], dtype=torch.long)
    line_tensor = lineToTensor(line)
    return category, line, category_tensor, line_tensor
```

```
def randomValidationExample():
    category = randomChoice(val_categories)
    line = randomChoice(val_category_lines[category])
    val_category_tensor = torch.tensor([val_categories.index(category)], dtype=torch.long)
    val_line_tensor = lineToTensor(line)
    return category, line, val_category_tensor, val_line_tensor
```

```
def shuffle_arrs(a,b,c,d):
    combined = list(zip(a, b, c, d))
    random.shuffle(combined)
    a, b, c, d = zip(*combined)
    return a,b,c,d
```

```
def genData(category_line_hash, categories_arr):
    x, y, x_tensor, y_tensor = [], [], [], []
    for y_category in category_line_hash.keys():
        for x_line in category_line_hash[y_category]:
            y.append(y_category)
            x.append(x_line)
            y_tensor.append(torch.tensor([categories_arr.index(y_category)], dtype=torch.long))
            x_tensor.append(lineToTensor(x_line))
    x, y, x_tensor, y_tensor = shuffle_arrs(x, y, x_tensor, y_tensor)
    return x, y, x_tensor, y_tensor
```

```
def TrainingData():
    return genData(category_lines, all_categories)
# v = 1
```

```

π y = []
# x = []
# for y_category in category_lines.keys():
#     for x_line in category_lines[y_category]:
#         y.append(y_category)
#         x.append(x_line)
#         y_tensor.append(torch.tensor([val_categories.index(category)], dtype=to
#         x_tensor.append(lineToTensor(x_line))
# x, y, x_tensor, y_tensor = shuffle_arrs(x, y, x_tensor, y_tensor)
# return x, y

def ValidationData():
    return genData(val_category_lines, val_categories)
# y = []
# x = []
# y_tensor = []
# x_tensor = []
# for y_category in val_category_lines.keys():
#     for x_line in val_category_lines[y_category]:
#         y.append(y_category)
#         x.append(x_line)
#         y_tensor.append(torch.tensor([val_categories.index(category)], dtype=to
#         x_tensor.append(lineToTensor(x_line))
# x, y, x_tensor, y_tensor = shuffle_arrs(x, y, x_tensor, y_tensor)
# return x, y

print("=== Train ===")
x,y,x_tensor,y_tensor= TrainingData()
print(x[:5])
print(y[:5])
# print(x_tensor[:1])
# print(y_tensor[:1])

print("=== Validation ===")
x,y,x_tensor,y_tensor = ValidationData()
print(x[:5])
print(y[:5])
# print(x_tensor[:1])
# print(y_tensor[:1])

[ ]> === Train ===
('vahimalaza', 'paihsing', 'arbab shafi muhammad', 'tak siah', 'pliringan')
('za', 'cn', 'af', 'ir', 'ir')
=== Validation ===
('fineview', 'magdebacken', 'la boussac', 'jafferabad', 'rexpoele')
('fi', 'de', 'fr', 'pk', 'fr')

```

▼ Training the Network

Now all it takes to train this network is show it a bunch of examples, have it make guesses, and tell it
 For the loss function `nn.NLLLoss` is appropriate, since the last layer of the RNN is `nn.LogSoftmax`.

```
criterion = nn.NLLLoss()
```

Each loop of training will:

- Create input and target tensors
- Create a zeroed initial hidden state
- Read each letter in and
 - Keep hidden state for next letter
- Compare final output to target
- Back-propagate
- Return the output and loss

```
learning_rate = 0.0001 # If you set this too high, it might explode. If too low, it m
```

```
def train(category_tensor, line_tensor):
    hidden = rnn.initHidden()

    rnn.zero_grad()

    # print("category_tensor={}, line_tensor.size()[0]={}".format(category_tensor, li
    for i in range(line_tensor.size()[0]):
        output, hidden = rnn(line_tensor[i], hidden)

    loss = criterion(output, category_tensor)
    loss.backward()

    # Add parameters' gradients to their values, multiplied by learning rate
    for p in rnn.parameters():
        p.data.add_(-learning_rate, p.grad.data)

    return output, loss.item()

# Just return an output given a line
def evaluate(line_tensor, category_tensor):
    hidden = rnn.initHidden()

    for i in range(line_tensor.size()[0]):
        output, hidden = rnn(line_tensor[i], hidden)
    loss = criterion(output, category_tensor)

    return output, loss.item()
```


Now we just have to run that with a bunch of examples. Since the `train` function returns both the output and the loss, we also keep track of loss for plotting. Since there are 1000s of examples we print only every `print_every` loss.

```
import time
import math

print_every = 1000 # total = 27000
plot_every = 1000 # 5000

# Keep track of losses for plotting
current_loss = 0
val_losses = 0.
train_acc_thru_time_aggregate, val_acc_thru_time_aggregate = 0., 0.
train_losses_thru_time = []
val_losses_thru_time = []
train_acc_thru_time = []
val_acc_thru_time = []

def timeSince(since):
    now = time.time()
    s = now - since
    m = math.floor(s / 60)
    s -= m * 60
    return '%dm %ds' % (m, s)

start = time.time()

print("learning rate = ", learning_rate)

x_train, y_train, x_train_tensor, y_train_tensor = TrainingData()
x_val, y_val, x_val_tensor, y_val_tensor = ValidationData()

x_train_len = len(x_train)
x_val_len = 10 # len(x_val)
print("x_train_len:", x_train_len, ", x_val_len:", x_val_len)

for i in range(x_train_len):
    # category, line, category_tensor, line_tensor = randomTrainingExample() # TODO:
    category = y_train[i]
    line = x_train[i]
    category_tensor = y_train_tensor[i]
    line_tensor = x_train_tensor[i]

    output, loss = train(category_tensor, line_tensor)
    current_loss += loss

    val_loss_per_train_data = 0
    val_correct_guess_count = 0
```

```

train_correct_guess_count = 0
# for j in range(x_val_len):
#     val_output, val_loss = evaluate(x_val_tensor[j], y_val_tensor[j])
#     val_loss_per_train_data += val_loss
for j in range(x_val_len):
    # Train accuracy calc
    train_category, _, train_category_tensor, train_line_tensor = randomTrainingE
    train_output, train_loss = evaluate(train_line_tensor, train_category_tensor)
    train_guess, _ = categoryFromOutput(train_output)
    train_correct_guess_count += int(train_guess == train_category)

    # Validation accuracy calc
    val_category, _, val_category_tensor, val_line_tensor = randomValidationExamp
    val_output, val_loss = evaluate(val_line_tensor, val_category_tensor)
    val_guess, _ = categoryFromOutput(val_output)
    val_correct_guess_count += int(val_guess == val_category)

    val_loss_per_train_data += val_loss

# Aggregate accuracy
train_acc_per_train_data = train_correct_guess_count / x_val_len
train_acc_thru_time_aggregate += train_acc_per_train_data
val_acc_per_train_data = val_correct_guess_count / x_val_len
val_acc_thru_time_aggregate += val_acc_per_train_data

# Aggregate validation loss
val_loss_per_train_data_ave = val_loss_per_train_data / x_val_len
val_losses += val_loss_per_train_data_ave

# Print iter number, loss, name and guess
if i % print_every == 0:
    print("iter = {}({:d}%) | time taken = {} | train_loss={:.4f}, val_loss(ave)=
    debug_x, debug_y, debug_x_tensor, debug_y_tensor = [], [], [], []

# Add current loss avg to list of losses
if i % plot_every == 0:
    train_losses_thru_time.append(current_loss / plot_every)
    val_losses_thru_time.append(val_losses / plot_every)
    current_loss = 0
    val_losses = 0

    print("iter = {}({:d}%) | time taken = {} | train_acc_thru_time_ave={}, val_a

    train_acc_thru_time.append(train_acc_thru_time_aggregate / plot_every)
    val_acc_thru_time.append(val_acc_thru_time_aggregate / plot_every)
    train_acc_thru_time_aggregate = 0
    val_acc_thru_time_aggregate = 0

```



```

learning rate = 0.0001
x_train_len: 27000 , x_val_len: 10
iter = 0(0%) | time taken = 0m 3s | train_loss=2.1915, val_loss(ave)=2.1812 | tra
iter = 0(0%) | time taken = 0m 3s | train_acc_thru_time_ave=0.0001, val_acc_thru
iter = 1000(3%) | time taken = 0m 36s | train_loss=2.2186, val_loss(ave)=2.1894
iter = 1000(3%) | time taken = 0m 36s | train_acc_thru_time_ave=0.093999999999999
iter = 2000(7%) | time taken = 1m 10s | train_loss=2.2477, val_loss(ave)=2.2309
iter = 2000(7%) | time taken = 1m 10s | train_acc_thru_time_ave=0.093799999999999
iter = 3000(11%) | time taken = 1m 43s | train_loss=2.1128, val_loss(ave)=2.2127
iter = 3000(11%) | time taken = 1m 43s | train_acc_thru_time_ave=0.097099999999999
iter = 4000(14%) | time taken = 2m 16s | train_loss=2.2728, val_loss(ave)=2.2318
iter = 4000(14%) | time taken = 2m 16s | train_acc_thru_time_ave=0.102199999999999
iter = 5000(18%) | time taken = 2m 49s | train_loss=2.1472, val_loss(ave)=2.2132
iter = 5000(18%) | time taken = 2m 49s | train_acc_thru_time_ave=0.097399999999999
iter = 6000(22%) | time taken = 3m 23s | train_loss=2.2164, val_loss(ave)=2.1952
iter = 6000(22%) | time taken = 3m 23s | train_acc_thru_time_ave=0.111199999999999
iter = 7000(25%) | time taken = 3m 56s | train_loss=2.1604, val_loss(ave)=2.1973
iter = 7000(25%) | time taken = 3m 56s | train_acc_thru_time_ave=0.119499999999999
iter = 8000(29%) | time taken = 4m 29s | train_loss=2.2503, val_loss(ave)=2.1864
iter = 8000(29%) | time taken = 4m 29s | train_acc_thru_time_ave=0.118399999999999
iter = 9000(33%) | time taken = 5m 2s | train_loss=2.1701, val_loss(ave)=2.2192
iter = 9000(33%) | time taken = 5m 2s | train_acc_thru_time_ave=0.124099999999999
iter = 10000(37%) | time taken = 5m 35s | train_loss=2.1631, val_loss(ave)=2.2131
iter = 10000(37%) | time taken = 5m 35s | train_acc_thru_time_ave=0.130199999999999
iter = 11000(40%) | time taken = 6m 8s | train_loss=2.1985, val_loss(ave)=2.1789
iter = 11000(40%) | time taken = 6m 8s | train_acc_thru_time_ave=0.132099999999999
iter = 12000(44%) | time taken = 6m 42s | train_loss=2.1330, val_loss(ave)=2.2091
iter = 12000(44%) | time taken = 6m 42s | train_acc_thru_time_ave=0.133399999999999
iter = 13000(48%) | time taken = 7m 15s | train_loss=2.3260, val_loss(ave)=2.1621
iter = 13000(48%) | time taken = 7m 15s | train_acc_thru_time_ave=0.135999999999999
iter = 14000(51%) | time taken = 7m 48s | train_loss=2.1949, val_loss(ave)=2.1926
iter = 14000(51%) | time taken = 7m 48s | train_acc_thru_time_ave=0.138199999999999
iter = 15000(55%) | time taken = 8m 21s | train_loss=2.1718, val_loss(ave)=2.2071
iter = 15000(55%) | time taken = 8m 21s | train_acc_thru_time_ave=0.144399999999999
iter = 16000(59%) | time taken = 8m 54s | train_loss=2.2873, val_loss(ave)=2.1996
iter = 16000(59%) | time taken = 8m 54s | train_acc_thru_time_ave=0.136699999999999
iter = 17000(62%) | time taken = 9m 27s | train_loss=2.2934, val_loss(ave)=2.2031
iter = 17000(62%) | time taken = 9m 27s | train_acc_thru_time_ave=0.139099999999999
iter = 18000(66%) | time taken = 10m 0s | train_loss=2.2057, val_loss(ave)=2.2011
iter = 18000(66%) | time taken = 10m 0s | train_acc_thru_time_ave=0.150599999999999
iter = 19000(70%) | time taken = 10m 33s | train_loss=2.1749, val_loss(ave)=2.2101
iter = 19000(70%) | time taken = 10m 33s | train_acc_thru_time_ave=0.149899999999999
iter = 20000(74%) | time taken = 11m 6s | train_loss=2.1967, val_loss(ave)=2.2291
iter = 20000(74%) | time taken = 11m 6s | train_acc_thru_time_ave=0.154899999999999
iter = 21000(77%) | time taken = 11m 39s | train_loss=2.1256, val_loss(ave)=2.1711
iter = 21000(77%) | time taken = 11m 39s | train_acc_thru_time_ave=0.161899999999999
iter = 22000(81%) | time taken = 12m 12s | train_loss=2.1767, val_loss(ave)=2.2111
iter = 22000(81%) | time taken = 12m 12s | train_acc_thru_time_ave=0.167699999999999
iter = 23000(85%) | time taken = 12m 45s | train_loss=2.1484, val_loss(ave)=2.1711
iter = 23000(85%) | time taken = 12m 45s | train_acc_thru_time_ave=0.165299999999999
iter = 24000(88%) | time taken = 13m 18s | train_loss=2.2119, val_loss(ave)=2.1911
iter = 24000(88%) | time taken = 13m 18s | train_acc_thru_time_ave=0.180899999999999
iter = 25000(92%) | time taken = 13m 50s | train_loss=2.1419, val_loss(ave)=2.1911
iter = 25000(92%) | time taken = 13m 50s | train_acc_thru_time_ave=0.177999999999999

```

```
import matplotlib.pyplot as plt
```

```
import matplotlib.ticker as ticker
```

```
plt.figure()
train_loss_plot = plt.plot(train_losses_thru_time[1:], label='Train Loss')
val_loss_plot = plt.plot(val_losses_thru_time[1:], label="Val Loss")
plt.legend()
```

```
print("train_losses_thru_time")
print(train_losses_thru_time[1:])
print("val_losses_thru_time")
print(val_losses_thru_time[1:])
```

```
☞ train_losses_thru_time
[2.1959270610809325, 2.1957004733085634, 2.194558796405792, 2.195638681650162, 2.
val_losses_thru_time
[2.2042345156431185, 2.2048589717865, 2.2042745894193683, 2.2041240722179407, 2.2
```



```
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
```

```
plt.figure()
train_acc_plot = plt.plot(train_acc_thru_time[1:], label='Train Accuracy')
val_acc_plot = plt.plot(val_acc_thru_time[1:], label="Val Accuracy")
plt.legend()
```


```
print("train_acc_thru_time")
print(train_acc_thru_time[1:])
print("val_acc_thru_time")
print(val_acc_thru_time[1:])
```

☞

```

train_acc_thru_time
[0.09399999999999994, 0.09379999999999977, 0.09709999999999974, 0.10219999999999999,
val_acc_thru_time
[0.09599999999999967, 0.09549999999999996, 0.09939999999999999, 0.10609999999999997,

```



```

import time
import math

print_every = 1000 # total = 27000
plot_every = 1000 # 5000

# Keep track of losses for plotting
current_loss = 0
val_losses = 0.
train_acc_thru_time_aggregate, val_acc_thru_time_aggregate = 0., 0.
train_losses_thru_time = []
val_losses_thru_time = []
train_acc_thru_time = []
val_acc_thru_time = []

def timeSince(since):
    now = time.time()
    s = now - since
    m = math.floor(s / 60)
    s -= m * 60
    return '%dm %ds' % (m, s)

start = time.time()

print("learning rate = ", learning_rate)

x_train, y_train, x_train_tensor, y_train_tensor = TrainingData()
x_val, y_val, x_val_tensor, y_val_tensor = ValidationData()

x_train_len = len(x_train)
x_val_len = 10 # len(x_val)

```

```

print("x_train_len:", x_train_len, ", x_val_len:", x_val_len)

for i in range(x_train_len):
    # category, line, category_tensor, line_tensor = randomTrainingExample() # TODO:
    category = y_train[i]
    line = x_train[i]
    category_tensor = y_train_tensor[i]
    line_tensor = x_train_tensor[i]

    output, loss = train(category_tensor, line_tensor)
    current_loss += loss

    val_loss_per_train_data = 0
    val_correct_guess_count = 0
    train_correct_guess_count = 0
    # for j in range(x_val_len):
    #     val_output, val_loss = evaluate(x_val_tensor[j], y_val_tensor[j])
    #     val_loss_per_train_data += val_loss
    for j in range(x_val_len):
        # Train accuracy calc
        train_category, _, train_category_tensor, train_line_tensor = randomTrainingE
        train_output, train_loss = evaluate(train_line_tensor, train_category_tensor)
        train_guess, _ = categoryFromOutput(train_output)
        train_correct_guess_count += int(train_guess == train_category)

        # Validation accuracy calc
        val_category, _, val_category_tensor, val_line_tensor = randomValidationExamp
        val_output, val_loss = evaluate(val_line_tensor, val_category_tensor)
        val_guess, _ = categoryFromOutput(val_output)
        val_correct_guess_count += int(val_guess == val_category)

    val_loss_per_train_data += val_loss

    # Aggregate accuracy
    train_acc_per_train_data = train_correct_guess_count / x_val_len
    train_acc_thru_time_aggregate += train_acc_per_train_data
    val_acc_per_train_data = val_correct_guess_count / x_val_len
    val_acc_thru_time_aggregate += val_acc_per_train_data

    # Aggregate validation loss
    val_loss_per_train_data_ave = val_loss_per_train_data / x_val_len
    val_losses += val_loss_per_train_data_ave

    # Print iter number, loss, name and guess
    if i % print_every == 0:
        print("iter = {}({:d}%) | time taken = {} | train_loss={:.4f}, val_loss(ave)=
        debug_x, debug_y, debug_x_tensor, debug_y_tensor = [], [], [], []

    # Add current loss avg to list of losses
    if i % plot_every == 0:
        train_losses_thru_time.append(current_loss / plot_every)
        val_losses_thru_time.append(val_losses / plot_every)

```

```
current_loss = 0
val_losses = 0

print("iter = {}({:d}%) | train_acc_thru_time_ave={}, val_acc_thru_time_ave={

train_acc_thru_time.append(train_acc_thru_time_aggregate / plot_every)
val_acc_thru_time.append(val_acc_thru_time_aggregate / plot_every)
train_acc_thru_time_aggregate = 0
val_acc_thru_time_aggregate = 0
```



```

learning rate = 0.0001
x_train_len: 27000 , x_val_len: 10
iter = 0(0%) | time taken = 0m 3s | train_loss=2.2156, val_loss(ave)=2.2050 | tra
iter = 0(0%) | train_acc_thru_time_ave=0m 3s, val_acc_thru_time_ave=0.0004
iter = 1000(3%) | time taken = 0m 37s | train_loss=2.1467, val_loss(ave)=2.2260
iter = 1000(3%) | train_acc_thru_time_ave=0m 37s, val_acc_thru_time_ave=0.194099
iter = 2000(7%) | time taken = 1m 10s | train_loss=2.2675, val_loss(ave)=2.1909
iter = 2000(7%) | train_acc_thru_time_ave=1m 10s, val_acc_thru_time_ave=0.200699
iter = 3000(11%) | time taken = 1m 43s | train_loss=2.1159, val_loss(ave)=2.2019
iter = 3000(11%) | train_acc_thru_time_ave=1m 43s, val_acc_thru_time_ave=0.20849
iter = 4000(14%) | time taken = 2m 17s | train_loss=2.1421, val_loss(ave)=2.1809
iter = 4000(14%) | train_acc_thru_time_ave=2m 17s, val_acc_thru_time_ave=0.20409
iter = 5000(18%) | time taken = 2m 50s | train_loss=2.1097, val_loss(ave)=2.2194
iter = 5000(18%) | train_acc_thru_time_ave=2m 50s, val_acc_thru_time_ave=0.20439
iter = 6000(22%) | time taken = 3m 24s | train_loss=2.1486, val_loss(ave)=2.2168
iter = 6000(22%) | train_acc_thru_time_ave=3m 24s, val_acc_thru_time_ave=0.21449
iter = 7000(25%) | time taken = 3m 57s | train_loss=2.3136, val_loss(ave)=2.1859
iter = 7000(25%) | train_acc_thru_time_ave=3m 57s, val_acc_thru_time_ave=0.21159
iter = 8000(29%) | time taken = 4m 30s | train_loss=2.1142, val_loss(ave)=2.1992
iter = 8000(29%) | train_acc_thru_time_ave=4m 30s, val_acc_thru_time_ave=0.21889
iter = 9000(33%) | time taken = 5m 3s | train_loss=2.2112, val_loss(ave)=2.1965
iter = 9000(33%) | train_acc_thru_time_ave=5m 3s, val_acc_thru_time_ave=0.211199
iter = 10000(37%) | time taken = 5m 36s | train_loss=2.2930, val_loss(ave)=2.2176
iter = 10000(37%) | train_acc_thru_time_ave=5m 36s, val_acc_thru_time_ave=0.21629
iter = 11000(40%) | time taken = 6m 9s | train_loss=2.1818, val_loss(ave)=2.2028
iter = 11000(40%) | train_acc_thru_time_ave=6m 9s, val_acc_thru_time_ave=0.22089
iter = 12000(44%) | time taken = 6m 43s | train_loss=2.2718, val_loss(ave)=2.1966
iter = 12000(44%) | train_acc_thru_time_ave=6m 43s, val_acc_thru_time_ave=0.22409
iter = 13000(48%) | time taken = 7m 16s | train_loss=2.1547, val_loss(ave)=2.1882
iter = 13000(48%) | train_acc_thru_time_ave=7m 16s, val_acc_thru_time_ave=0.22329
iter = 14000(51%) | time taken = 7m 49s | train_loss=2.1200, val_loss(ave)=2.1981
iter = 14000(51%) | train_acc_thru_time_ave=7m 49s, val_acc_thru_time_ave=0.22799
iter = 15000(55%) | time taken = 8m 22s | train_loss=2.1782, val_loss(ave)=2.2247
iter = 15000(55%) | train_acc_thru_time_ave=8m 22s, val_acc_thru_time_ave=0.23799
iter = 16000(59%) | time taken = 8m 55s | train_loss=2.1858, val_loss(ave)=2.1932
iter = 16000(59%) | train_acc_thru_time_ave=8m 55s, val_acc_thru_time_ave=0.23389
iter = 17000(62%) | time taken = 9m 28s | train_loss=2.1522, val_loss(ave)=2.2064
iter = 17000(62%) | train_acc_thru_time_ave=9m 28s, val_acc_thru_time_ave=0.23599
iter = 18000(66%) | time taken = 10m 1s | train_loss=2.2886, val_loss(ave)=2.2126
iter = 18000(66%) | train_acc_thru_time_ave=10m 1s, val_acc_thru_time_ave=0.23929
iter = 19000(70%) | time taken = 10m 34s | train_loss=2.2778, val_loss(ave)=2.2071
iter = 19000(70%) | train_acc_thru_time_ave=10m 34s, val_acc_thru_time_ave=0.2381
iter = 20000(74%) | time taken = 11m 7s | train_loss=2.2315, val_loss(ave)=2.2306
iter = 20000(74%) | train_acc_thru_time_ave=11m 7s, val_acc_thru_time_ave=0.23029
iter = 21000(77%) | time taken = 11m 41s | train_loss=2.2578, val_loss(ave)=2.2411
iter = 21000(77%) | train_acc_thru_time_ave=11m 41s, val_acc_thru_time_ave=0.2416
iter = 22000(81%) | time taken = 12m 14s | train_loss=2.1304, val_loss(ave)=2.2112
iter = 22000(81%) | train_acc_thru_time_ave=12m 14s, val_acc_thru_time_ave=0.2477
iter = 23000(85%) | time taken = 12m 48s | train_loss=2.1302, val_loss(ave)=2.1851
iter = 23000(85%) | train_acc_thru_time_ave=12m 48s, val_acc_thru_time_ave=0.2412
iter = 24000(88%) | time taken = 13m 21s | train_loss=2.1176, val_loss(ave)=2.1971
iter = 24000(88%) | train_acc_thru_time_ave=13m 21s, val_acc_thru_time_ave=0.2491
iter = 25000(92%) | time taken = 13m 54s | train_loss=2.1380, val_loss(ave)=2.2321
iter = 25000(92%) | train_acc_thru_time_ave=13m 54s, val_acc_thru_time_ave=0.2497

```


▼ Plotting the Results


Plotting the historical loss from `all_losses` shows the network learning:

```
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
```

```
plt.figure()
train_loss_plot = plt.plot(train_losses_thru_time[1:], label='Train Loss')
val_loss_plot = plt.plot(val_losses_thru_time[1:], label="Val Loss")
plt.legend()
```

```
print("train_losses_thru_time")
print(train_losses_thru_time[1:])
print("val_losses_thru_time")
print(val_losses_thru_time[1:])
```

```
☞ train_losses_thru_time
[2.1852655036449433, 2.181572945356369, 2.180067013978958, 2.1807363028526305, 2.
val_losses_thru_time
[2.205783145713804, 2.2065187231540686, 2.205634904360772, 2.2068484337568304, 2.
```



```
train_acc_thru_time_aggregate
```

```
☞ 250.09999999999943
```

```
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
```

```
plt.figure()
train_acc_plot = plt.plot(train_acc_thru_time[1:], label='Train Accuracy')
val_acc_plot = plt.plot(val_acc_thru_time[1:], label="Val Accuracy")
plt.legend()
```


```
print("train_acc_thru_time")
```

```
print(train_acc_thru_time[1:])
print("val_acc_thru_time")
print(val_acc_thru_time[1:])
```

```

[ ]> train_acc_thru_time
[0.19409999999999897, 0.20069999999999902, 0.20849999999999894, 0.20409999999999998]
val_acc_thru_time
[0.18379999999999905, 0.19479999999999914, 0.18999999999999928, 0.20319999999999998]

```



▼ Evaluating the Results

To see how well the network performs on different categories, we will create a confusion matrix, indicating which language the network guesses (columns). To calculate the confusion matrix a bunch of samples, we use `evaluate()`, which is the same as `train()` minus the backprop.

```
# Keep track of correct guesses in a confusion matrix
confusion = torch.zeros(n_categories, n_categories)
confusion_no_norm = torch.zeros(n_categories, n_categories)

n_confusion = 500 # 10000

# Go through a bunch of examples and record which are correctly guessed
for i in range(n_confusion):
    category, line, category_tensor, line_tensor = randomTrainingExample()
    output, _ = evaluate(line_tensor, category_tensor)
    guess, guess_i = categoryFromOutput(output)
    category_i = all_categories.index(category)
    confusion[category_i][guess_i] += 1
    confusion_no_norm[category_i][guess_i] += 1

# Normalize by dividing every row by its sum
for i in range(n_categories):
    confusion[i] = confusion[i] / confusion[i].sum()

# Set up plot
fig = plt.figure()
```

```


fig = plt.figure()
ax = fig.add_subplot(111)
cax = ax.matshow(confusion.numpy())
fig.colorbar(cax)

# Set up axes
ax.set_xticklabels([''] + all_categories, rotation=90)
ax.set_yticklabels([''] + all_categories)

# Force label at every tick
ax.xaxis.set_major_locator(ticker.MultipleLocator(1))
ax.yaxis.set_major_locator(ticker.MultipleLocator(1))

# sphinx_gallery_thumbnail_number = 2
plt.show()

```



```

print("confusion matrix (no normalization)")
print(confusion_no_norm)

def get_multi_class_accuracy(confusion):
    total = torch.sum(confusion)
    correct = 0
    for guess_i in range(len(confusion)):
        correct += confusion[guess_i][guess_i]
    return correct / total

print("accuracy, multi-class = {}".format(get_multi_class_accuracy(confusion_no_norm)

```

☐→

```

        confusion matrix (no normalization)
def get_pos_tp(target_i, confusion):
    pos = torch.sum(confusion[:, target_i])
    tp = confusion[target_i][target_i]
    return pos, tp

def get_multi_class_precision(confusion):
    '''
    multi-class-precision = sum(all tp's across class) / sum(all pos' across class)
    '''
    pos = 0
    tp = 0
    for i in range(len(confusion)):
        target_pos, target_tp = get_pos_tp(i, confusion)
        pos += target_pos
        tp += target_tp

    precision = tp / pos
    return precision

print("precision, multi-class = {}".format(get_multi_class_precision(confusion_no_nor

↳ precision, multi-class = 0.23600000143051147

```

You can pick out bright spots off the main axis that show which languages it guesses incorrectly, e.g. Italian. It seems to do very well with Greek, and very poorly with English (perhaps because of overlap

▼ Running on User Input

```

# def predict(input_line, n_predictions=3):
#     print('\n> %s' % input_line)
#     with torch.no_grad():
#         output = evaluate(lineToTensor(input_line))

#         # Get top N categories
#         topv, topi = output.topk(n_predictions, 1, True)
#         predictions = []

#         for i in range(n_predictions):
#             value = topv[0][i].item()
#             category_index = topi[0][i].item()
#             print('(%2f) %s' % (value, all_categories[category_index]))
#             predictions.append([value, all_categories[category_index]])

# predict('Dovesky')
# predict('Jackson')
# predict('Satoshi')

```

The final versions of the scripts in the Practical PyTorch repo <<https://github.com/spro/p-rnn-classification>>__ split the above code into a few files:

- `data.py` (loads files)
- `model.py` (defines the RNN)
- `train.py` (runs training)
- `predict.py` (runs `predict()` with command line arguments)
- `server.py` (serve prediction as a JSON API with bottle.py)

Run `train.py` to train and save the network.

Run `predict.py` with a name to view predictions:

::

```
$ python predict.py Hazaki
(-0.42) Japanese
(-1.39) Polish
(-3.51) Czech
```

Run `server.py` and visit <http://localhost:5533/Yourname> to get JSON output of predictions.

Exercises

- Try with a different dataset of line -> category, for example:
 - Any word -> language
 - First name -> gender
 - Character name -> writer
 - Page title -> blog or subreddit
- Get better results with a bigger and/or better shaped network
 - Add more linear layers
 - Try the `nn.LSTM` and `nn.GRU` layers
 - Combine multiple of these RNNs as a higher level network

